

# OPERATING SYSTEMS

**Subject code : CSC 404**



**Subject In-charge**

**Nidhi Gaur**

**Assistant Professor**


email: [nidhigaur@sfit.ac.in](mailto:nidhigaur@sfit.ac.in)



# Module 3:Deadlocks



# Bridge Crossing Example



- Traffic only in one direction.
- Each section of a bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources ).
- Several cars may have to be backed up if a deadlock occurs.
- Starvation is possible.



# Deadlock

---

- Permanent blocking of a set of processes that either compete for system resources or communicate with each other
- No efficient solution
- Involve conflicting needs for resources by two or more processes




# System Model


---

- Various types of resources in multiprogramming environment.
- Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances.
- Each process utilizes a resource as follows:
  - request
  - use
  - release





(a) Deadlock possible



(b) Deadlock

**Figure 6.1 Illustration of Deadlock**

# Deadlock

A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.

## Example

- System has 2 disk drives.
- $P_1$  and  $P_2$  each hold one disk drive and each needs another one.

## Example

- semaphores  $A$  and  $B$ , initialized to 1

$P_0$   
wait (A);  
wait (B);

$P_1$   
wait(B)  
wait(A)



# Consumable Resources

---

- Created (produced) and destroyed (consumed) by a process
- Interrupts, signals, messages, and information in I/O buffers
- May take a rare combination of events to cause deadlock





# Reusable Resources

---

- Used by one process at a time and not depleted by that use
- Processes obtain resources that they later release for reuse by other processes
- Processors, I/O channels, main and secondary memory, files, databases, and semaphores
- Deadlock occurs if each process holds one resource and requests the other



# Another Example

- Space is available for allocation of 200K bytes, and the following sequence of events occur
- Deadlock occurs if both processes progress to their second request



# Deadlock Characterization

---


**Deadlock can arise if four conditions hold simultaneously.**

- **Mutual exclusion:** only one process at a time can use a resource.
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.



# Deadlock Characterization

- **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .



# Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .

$V$  is partitioned into two types:

–  $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system.

–  $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system.


**Request edge** – directed edge  $P_1 \rightarrow R_j$

**Assignment edge** – directed edge  $R_j \rightarrow P_i$




# Resource-Allocation Graph (Cont.)



Process



Resource Type with 4 instances





$P_i$  requests instance of  $R_j$



$R_j$


$P_i$  is holding an instance of  $R_j$




$R_j$



# Example of a Resource Allocation Graph




# Resource Allocation Graph With A Deadlock





# Graph With A Cycle But No Deadlock



# Basic Facts

---

- If graph contains no cycles  $\Rightarrow$  no deadlock.
- If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock.
  - if several instances per resource type, possibility of deadlock.



# Deadlock Prevention

---

Restrain the ways request can be made.

- **Mutual Exclusion** – must hold for non sharable resources.
- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources.

Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none.



# Deadlock Prevention (Cont.)

## ➤ No Preemption –

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
- Preempted resources are added to the list of resources for which the process is waiting.
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.

## ➤ Circular wait

- impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration.



# Deadlock Avoidance

---

- A decision is made dynamically whether the current resource allocation request will, if granted, potentially lead to a deadlock
- Requires knowledge of future process request



# Deadlock avoidance

---

- What is the difference between deadlock prevention and deadlock avoidance?



# Deadlock Avoidance

---

- Deadlock avoidance is concerned with starting with an environment where deadlock is theoretically possible, but by some algo in the OS , it is ensured, before allocating any resource that after allocating it, a deadlock can be avoided.
- Deadlock prevention was concerned with imposing certain restrictions on the environment or processes so that deadlocks can never occur.



# Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state.
- System is in **safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the systems such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$ .

That is:

- If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished.






- **State** of the system is the current allocation of resources to process
- **Safe state** is where there is at least one sequence that does not result in deadlock
- **Unsafe state** is a state that is not safe, can lead to deadlock.



# Safe, Unsafe, Deadlock State



# Basic Facts

---

- If a system is in safe state  $\Rightarrow$  no deadlocks.
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock.
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.



# Avoidance algorithms


---

- Single instance of a resource type.  
Use a resource-allocation graph
- Multiple instances of a resource type. Use the banker's algorithm





# Resource-Allocation Graph Scheme

- **Claim edge**  $P_i \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$ ; represented by a dashed line.
- Claim edge converts to request edge when a process requests a resource.
- Request edge converts to an assignment edge when the resource is allocated to the process.
- When a resource is released by a process, assignment edge reconverts to a claim edge.
- Resources must be claimed *a priori* in the system.



# Unsafe State In Resource-Allocation

## Graph



# Resource-Allocation Graph

## Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if **converting the request edge to an assignment edge does not result in the formation of a cycle** in the resource allocation graph



# Banker's Algorithm

---

- Multiple instances.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.





# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

**Available:** Vector of length  $m$ . If available  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.

**Max:**  $n \times m$  matrix. If  $Max[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .

**Allocation:**  $n \times m$  matrix. If  $Allocation[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .

**Need:**  $n \times m$  matrix. If  $Need[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.

$$Need[i,j] = Max[i,j] - Allocation[i,j].$$



# Safety Algorithm

1. Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively. Initialize:

*Work* = *Available*

*Finish* [ $i$ ] = *false* for  $i = 0, 1, \dots, n-1$ .

2. Find an  $i$  such that both:

(a) *Finish* [ $i$ ] = *false*

*// process i not yet done*

(b)  $Need_i \leq Work$

*// its need can be satisfied*

If no such  $i$  exists, go to step 4.

3. *Work* = *Work* + *Allocation* <sub>$i$</sub>

*// run it and reclaim*

*Finish* [ $i$ ] = *true*

*// process i completes*

go to step 2.

4. If *Finish* [ $i$ ] == *true* for all  $i$ , then the system is in a safe state.



# Example

- Consider a system with 12 magnetic tape drives. Snapshot at time T0 is:

Process	Maximum needs	Allocated
P0	10	5
P1	4	2
P2	9	2

Available
3

**Find need matrix and safe sequence.**



# NEED Matrix

Process	Maximum needs	Allocated	Need = Max- Allocation
P0	10	5	5
P1	4	2	2
P2	9	2	7



# Example

- Consider a system with 12 magnetic tape drives. Snapshot at time T0 is:

Process	Maximum needs	Allocated	Available/ work	Need = Max- Allocation
P0	10	5	3	5
P1	4	2		2
P2	9	2		7

1) Let Work vector be length m  
Finish be vector of length n  
Initialize Work = Available  
Finish[i] = false, for i = 0,1,2,...,n-1

2) Find an index i such that

- Finish[i] == false
- Need i <= work

If no such i exists, go to step 4.

3) Work = work + Allocation  
Finish[i] = true  
Go to step 2.

4) If Finish[i] == true for all i, then system is in safe state.

1 P0 Need <= Work **NO**

2 P1 Need <= Work **YES**

Work = Work + Allocation  
= 3+2 =5

3. P2 Need <= Work **NO**

4. P0 Need <= Work **YES**

Work = Work + Allocation  
= 5+5=10

5. P2 Need <= Work **YES**

Work = Work + Allocation  
= 10+2=12

# Example

---

- Safe sequence  $\langle P1 \ P0 \ P2 \rangle$



# Example of Banker's Algorithm

5 processes  $P_0$  through  $P_4$ ;

3 resource types:

A (10 instances), B (5 instances), and C (7 instances).

Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

Find the need matrix and safe sequence.



# Example (Cont.)

---

The content of the matrix *Need* is defined to be *Max – Allocation*.

	<u><i>Need</i></u>		
	<i>A</i>	<i>B</i>	<i>C</i>
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1





# Example of Banker's Algorithm

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>	<u>Need</u>
	A B C	A B C	A B C	A B C
$P_0$	0 1 0	7 5 3	3 3 2	
$P_1$	2 0 0	3 2 2		7 4 3
$P_2$	3 0 2	9 0 2		1 2 2
$P_3$	2 1 1	2 2 2		6 0 0
$P_4$	0 0 2	4 3 3		0 1 1
				4 3 1

Total resources = Allocation + Available

$$\begin{array}{r}
 \text{A B C} \\
 = \quad 10 \ 5 \ 7
 \end{array}$$



# Example of Banker's algo(safety)

1. P0 : Need <= Work

$$7\ 4\ 3 \leq 3\ 3\ 2 \quad \text{NO}$$

2. P1 : Need <= Work

$$1\ 2\ 2 \leq 3\ 3\ 2 \quad \text{YES}$$

Work = Work + allocation

$$3\ 3\ 2 + 2\ 0\ 0$$

$$= 5\ 3\ 2$$

3 P2: Need <= Work

$$6\ 0\ 0 \leq 5\ 3\ 2 \quad \text{NO}$$

3 P3: Need <= Work YES

$$0\ 1\ 1 \leq 5\ 3\ 2$$

$$5\ 3\ 2 + 2\ 1\ 1 = 7\ 4\ 3$$

5. P4: Need <= Work YES

$$4\ 3\ 1 \leq 7\ 4\ 3$$

Work = Work + allocation

$$7\ 4\ 3 + 0\ 0\ 2 = 7\ 4\ 5$$

6. P0: Need <= Work

$$7\ 4\ 3 \leq 7\ 4\ 5 \quad \text{YES}$$

Work = 7 4 5 + 0 1 0

$$= 7\ 5\ 5$$

7. P2 : Need <= Work YES

Work = 7 5 5 + 3 0 2

$$= 10\ 5\ 7$$

Available/ WORK

A B C

3 3 2

Need

A B C

P0 7 4 3

P1 1 2 2

P2 6 0 0

P3 0 1 1

P4 4 3 1



# Example (Cont.)

The content of the matrix *Need* is defined to be *Max – Allocation*.

	<u><i>Need</i></u>
	<i>A B C</i>
$P_0$	7 4 3
$P_1$	1 2 2
$P_2$	6 0 0
$P_3$	0 1 1
$P_4$	4 3 1

The system is in a safe state since the sequence  
 $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety criteria.



# Example

	MAX			ALLOCATION		
PROCESS	R1	R2	R3	R1	R2	R3
P1	5	6	3	2	1	0
P2	8	5	6	3	2	3
P3	4	8	2	3	0	2
P4	7	4	3	3	2	0
P5	4	3	3	1	0	1

TOTAL Res		
R1	R2	R3
15	8	8

Q1. Find out available resources  
 Q2. Calculate need matrix.  
 Q2. Determine whether system is in safe state or no?



# example

- Available = total – allocation  
= 3 3 2

SAFE SEQUENCE:  
< P5 P4 P1 P2 P3 >

Need matrix

R1	R2	R3
3	5	3
5	3	3
1	9	0
4	2	3
3	3	2



# Resource-Request Algorithm for Process $P_i$

*Request* = request vector for process  $P_i$ . If  $Request_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$ .

1. If  $Request_i \leq Need_i$  go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If  $Request_i \leq Available$ , go to step 3. Otherwise  $P_i$  must wait, since resources are not available.
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- *If safe  $\Rightarrow$  the resources are allocated to  $P_i$ .*
- *If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored*



# Example: Resource Request

Check that Request  $\leq$  Available ; that is say P1 requests  
 $(1,0,2) \leq (3,3,2) \Rightarrow \text{true.}$

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	3 3 2
$P_1$	3 0 2	1 2 2	
$P_2$	3 0 1	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

What about request (3, 3,0) by P1 ?



# Deadlock Detection

---

In **deadlock detection**, requested resources are granted to processes whenever possible. Periodically the OS performs an algorithm that allows it to detect the circular wait condition.

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme





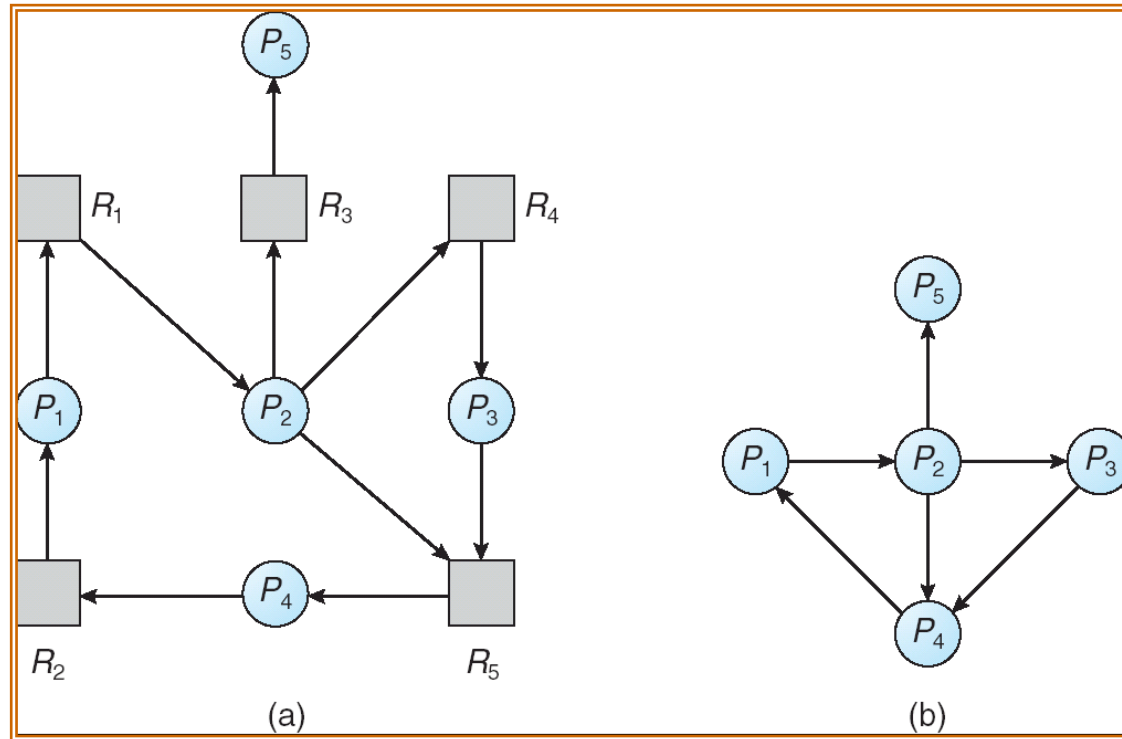
# Single Instance of Each Resource Type

---

- Maintain *wait-for* graph
  - Nodes are processes.
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$ .
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock.
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph.



# Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph



# Several Instances of a Resource Type

---

**Available:** A vector of length  $m$  indicates the number of available resources of each type.

**Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.

**Request:** An  $n \times m$  matrix indicates the current request of each process. If  $Request[i,j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .



# Deadlock Detection

## ALGORITHM

1. Mark each process that has a row in the Allocation matrix of all zeros.
  2. Initialize a temporary vector  $W$  to equal the available vector.
  3. Find an index  $i$  such that process  $i$  is currently unmarked and the  $i$ th row of  $Q$  (request matrix) is less than or equal to  $W$  i.e.  $Q_{ik} \leq W_k$ . If no such row is found terminate.
  4. If such a row is found, mark process  $i$  and add the row of allocation matrix to  $W$  i.e. set  $W_k = W_k + A_{ik}$
- return to step 3.



# Deadlock Detection

	R1	R2	R3	R4	R5
P1	0	1	0	0	1
P2	0	0	1	0	1
P3	0	0	0	0	1
P4	1	0	1	0	1

Request Matrix Q

	R1	R2	R3	R4	R5
P1	1	0	1	1	0
P2	1	1	0	0	0
P3	0	0	0	1	0
P4	0	0	0	0	0

Allocation Matrix A

R1	R2	R3	R4	R5
2	1	1	2	1

Resource Vector

R1	R2	R3	R4	R5
0	0	0	0	1

Available Vector

## ALGORITHM

1. Mark each process that has a row in the Allocation matrix of all zeros.
2. Initialize a temporary vector W to equal the available vector.
3. Find an index i such that process i is currently unmarked and the ith row of Q(request matrix) is less than or equal to W i.e.  $Q_{ik} \leq W_k$ . If no such row is found terminate.
4. If such a row is found, mark process i and add the row of allocation matrix

Example deadlock detection

# Deadlock Detection

- Mark P4 as P4 has no allocated resources.
- Set  $W = (00001)$
- Requirement of Process P3 is  $\leq W$ , so mark P3 and set  $W = W + (00010)$
- Hence  $W = (00011)$
- No other unmarked process has a row in  $Q$  that is less than or equal to  $W$ .

Therefore terminate the algorithm.

P1 and P2 unmarked indicates these are deadlocked.



# Strategies once Deadlock Detected

- Abort all deadlocked processes
- **Back up each deadlocked process** to some previously defined checkpoint, and restart all process
- **Successively abort deadlocked processes** until deadlock no longer exists
- **Successively preempt resources** until deadlock no longer exists



# Selection Criteria Deadlocked Processes

- Least amount of processor time consumed so far
- Least number of lines of output produced so far
- Most estimated time remaining
- Least total resources allocated so far
- Lowest priority





# Detection-Algorithm Usage

---

When, and how often, to invoke depends on:

- How often a deadlock is likely to occur?
- How many processes will need to be rolled back?
  - one for each disjoint cycle

If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock.



# Recovery from Deadlock: Process Termination

---

Abort all deadlocked processes.

Abort one process at a time until the deadlock cycle is eliminated.

In which order should we choose to abort?

- Priority of the process.
- How long process has computed, and how much longer to completion.
- Resources the process has used.
- Resources process needs to complete.
- How many processes will need to be terminated.
- Is process interactive or batch?




# Recovery from Deadlock: Resource Preemption

- Selecting a victim – minimize cost.
- Rollback – return to some safe state, restart process for that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.




# Dining Philosophers Problem



# Dining Philosophers Problem


Consider five philosophers who spend their lives thinking and eating.

- Philosophers share a common circular table surrounded by five chairs, each belonging to one philosopher.
- In the center of the table is a bowl of noodles/rice, and the table is laid with 5 plates and 5 chopsticks/spoons
- When philosopher is hungry he picks up two chopsticks that are closest, a philosopher may pick only one chopstick at a time.




# Dining Philosophers Problem

- When a hungry philosopher has both the chopsticks/spoons at the same time, he eats without releasing his chopsticks/spoons.
- When he is finished eating, he puts down both of his chopsticks/spoons and starts thinking again.



# Dining Philosophers Problem


- This is a classic problem to understand the synchronization among processes in an environment where there is scarcity of resources.
- To eat they must use two chopsticks or two spoons but there is scarcity of chopsticks or spoons.



# Dining Philosophers Problem

The material in this presentation belongs to St. Francis Institute of Technology and is solely for educational purposes. Distribution and modifications of the content is prohibited.


- The two conditions that need to be noted are:
  1. There are only five spoons available, each placed between two philosophers
  2. Philosophers can start eating only when they have two spoons. They can't eat with one spoon. It means that every philosopher takes one spoon from left and another from right and start eating. After eating put down both the spoons.





# Dining Philosophers Problem

- The problem here is obvious that all the philosophers cannot eat together as there are only five spoons whereas there is demand for 10.
- The solution should be such that if one philosopher is eating with one spoon, the adjacent philosopher should wait even if he is able to get one spoon.
- If all philosophers pick up one spoon at the same time, then this leads to deadlock, since every philosopher is waiting for other to release and no one is releasing.
- So a philosopher should not hold a spoon if he is not able to get the other one.




# Dining Philosophers Problem

- One simple solution is to represent each spoon by a semaphore, a spoon is mutually exclusive so we consider semaphore to protect each spoon.
- A philosopher tries to grab the spoon by executing a wait operation on that semaphore; he releases his spoon by executing the signal on the appropriate semaphore.
- We take an array of five elements


Shared data: semaphore spoon[5];

Where all elements of spoon are initialized to 1.




# Dining Philosophers Problem

```
semaphore spoon[5]=[1];  
int i;  
void philosopher (int i)  
{ while (true)  
{ think();  
wait(spoon[i]);  
wait(spoon[(i+1) mod 5];  
eat();  
signal(spoon[(i+1)mod 5];  
signal(spoon[i]);  
}  
}
```



What if all philosophers put wait on their left spoon first then can enter into deadlock



# Dining Philosophers Problem

The material in this presentation belongs to St. Francis Institute of Technology and is not to be used for educational purposes. Distribution and modifications of the content is prohibited.

Solution to the problem that is free of deadlocks:-

- Allow at most four philosophers to be sitting simultaneously at the table.
- Allow a philosopher to pick up spoons only if both spoons are available(must pick them in critical section.)

